

# tmdb-notebook

January 11, 2023

## 1 Project: TMDB MOVIES DATASET

TMDB dataset has more than 10000 movies has been collected from the movie database (TMDB), including features like movie title, budget, revenue, cast, director, runtime, genre ...etc, in this project, we will analyze this dataset to get some useful insights, but first we need to clean the data like removing irrelevant data to our analysis, dealing with missing data and correcting the bad values like zero budget or revenue or runtime.. etc

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import requests
import json
%matplotlib inline
```

```
[ ]: # importing our csv file to make a dataframe
df = pd.read_csv('tmdb-movies.csv')
df.head(3)
```

```
[ ]:      id    imdb_id  popularity    budget    revenue  original_title \
0  135397  tt0369610   32.985763  150000000  1513528810      Jurassic World
1    76341  tt1392190   28.419936  150000000   378436354  Mad Max: Fury Road
2   262500  tt2908446   13.112507  110000000   295238201      Insurgent
```

```
                                cast \
0  Chris Pratt|Bryce Dallas Howard|Irrfan Khan|Vi...
1  Tom Hardy|Charlize Theron|Hugh Keays-Byrne|Nic...
2  Shailene Woodley|Theo James|Kate Winslet|Ansel...
```

```
                                homepage    director \
0                                http://www.jurassicworld.com/  Colin Trevorrow
1                                http://www.madmaxmovie.com/    George Miller
2  http://www.thedivergentseries.movie/#insurgent  Robert Schwentke
```

```
                                tagline ... \
0                                The park is open. ...
1                                What a Lovely Day. ...
```

2 One Choice Can Destroy You ...

```
                                overview runtime \
0  Twenty-two years after the events of Jurassic ...      124
1  An apocalyptic story set in the furthest reach...      120
2  Beatrice Prior must confront her inner demons ...      119
```

```
                                genres \
0  Action|Adventure|Science Fiction|Thriller
1  Action|Adventure|Science Fiction|Thriller
2  Adventure|Science Fiction|Thriller
```

```
                                production_companies release_date vote_count \
0  Universal Studios|Amblin Entertainment|Legenda...      6/9/15      5562
1  Village Roadshow Pictures|Kennedy Miller Produ...      5/13/15      6185
2  Summit Entertainment|Mandeville Films|Red Wago...      3/18/15      2480
```

```
                                vote_average release_year budget_adj revenue_adj
0                                6.5          2015  1.379999e+08  1.392446e+09
1                                7.1          2015  1.379999e+08  3.481613e+08
2                                6.3          2015  1.012000e+08  2.716190e+08
```

[3 rows x 21 columns]

Next, we need to answer the next questions: Q1: how the number of produced movies changes over the years? Q2: Top 10 directors have the most revenue for thier films? Q3: what is the top 10 movies in revenue? Q4: what is the most produced genres ? Q5: does popularity affect the revenue? Q6: does runtime affect the revenue? Q7: does vote count affect the revenue? Q8: does vote average affect the revenue?

# Data Wrangling

### 1.0.1 \* General Properties

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10866 non-null  int64
1   imdb_id              10856 non-null  object
2   popularity            10866 non-null  float64
3   budget               10866 non-null  int64
4   revenue              10866 non-null  int64
5   original_title       10866 non-null  object
```

```

6   cast                10790 non-null object
7   homepage            2936 non-null object
8   director            10822 non-null object
9   tagline             8042 non-null object
10  keywords             9373 non-null object
11  overview            10862 non-null object
12  runtime             10866 non-null int64
13  genres              10843 non-null object
14  production_companies 9836 non-null object
15  release_date        10866 non-null object
16  vote_count          10866 non-null int64
17  vote_average        10866 non-null float64
18  release_year        10866 non-null int64
19  budget_adj          10866 non-null float64
20  revenue_adj         10866 non-null float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB

```

```
[ ]: print('duplicated values count:', df.duplicated().sum())
```

```
duplicated values count: 1
```

```
[ ]: df.describe()
```

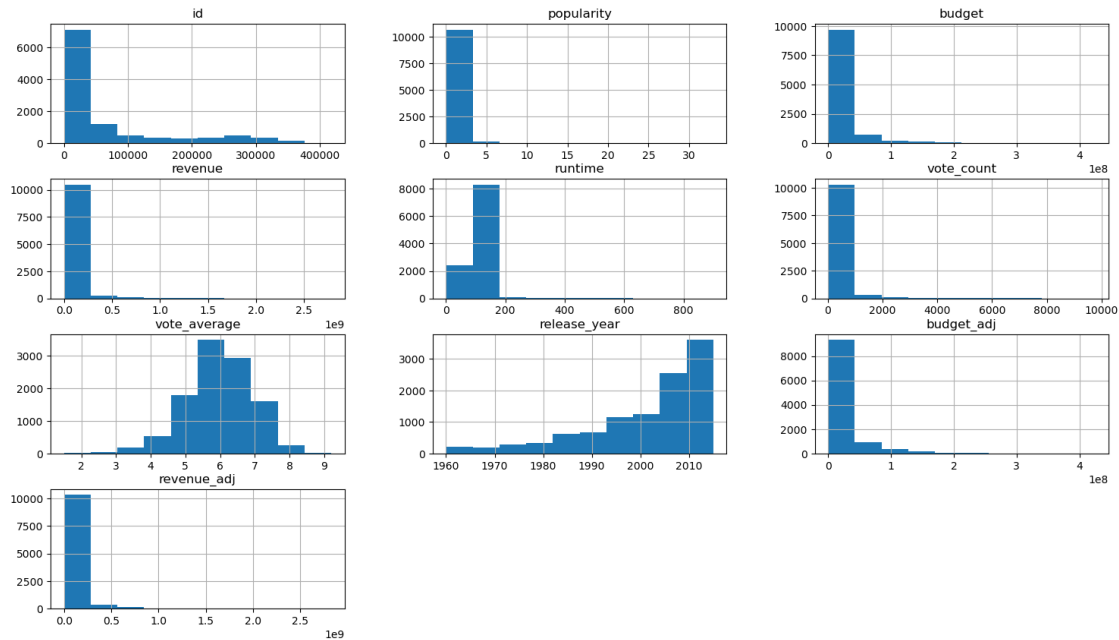
```
[ ]:
```

	id	popularity	budget	revenue	runtime \
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04	10866.000000
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07	102.070863
std	92130.136561	1.000185	3.091321e+07	1.170035e+08	31.381405
min	5.000000	0.000065	0.000000e+00	0.000000e+00	0.000000
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00	90.000000
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00	99.000000
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07	111.000000
max	417859.000000	32.985763	4.250000e+08	2.781506e+09	900.000000

	vote_count	vote_average	release_year	budget_adj	revenue_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07	5.136436e+07
std	575.619058	0.935142	12.812941	3.430616e+07	1.446325e+08
min	10.000000	1.500000	1960.000000	0.000000e+00	0.000000e+00
25%	17.000000	5.400000	1995.000000	0.000000e+00	0.000000e+00
50%	38.000000	6.000000	2006.000000	0.000000e+00	0.000000e+00
75%	145.750000	6.600000	2011.000000	2.085325e+07	3.369710e+07
max	9767.000000	9.200000	2015.000000	4.250000e+08	2.827124e+09

```
[ ]: df.hist(figsize=(18, 10));
```



from describe function we can see there is movies has ZERO VALUES for budget, revenue And runtime, lets see how many rows has zeros

```
[ ]: print('zero values count in budget column is: ', (df['budget'] == 0).sum())
      print('zero values count in revenue column is: ', (df['revenue'] == 0).sum())
      print('zero values count in runtime column is: ', (df['runtime'] == 0).sum())
```

```
zero values count in budget column is: 5696
zero values count in revenue column is: 6016
zero values count in runtime column is: 31
```

### 1.0.2 \* Data Cleaning

\* first we will drop the irrelated columns, and duplicated rows \* then will deal with the missing values in genres and director columns using tmdb official api to get the actual values \* also we will use the api to replace zeros in the budget, revenue and runtime columns.

```
[ ]: # dropping irrelated columns
df.drop(
    columns=[
        "homepage",
        "release_date",
        "tagline",
        "production_companies",
        "keywords",
        "overview",
        "imdb_id",
```

```

        "budget_adj",
        "revenue_adj",
        "cast",
    ],
    inplace=True,
)

```

```
[ ]: # converting all string columns to lower case
```

```

df.director, df.original_title, df.genres = (
    df.director.str.lower(),
    df.original_title.str.lower(),
    df.genres.str.lower(),
)

```

```
[ ]: # we have only 1 duplicated row, so can drop it
```

```

df.drop_duplicates(inplace=True)
# check again for duplicates
df.duplicated().sum()

```

```
[ ]: 0
```

```
[ ]: # lets check for null values
```

```
df.isnull().sum()
```

```

[ ]: id                0
    popularity         0
    budget             0
    revenue            0
    original_title     0
    director           44
    runtime            0
    genres             23
    vote_count         0
    vote_average       0
    release_year       0
    dtype: int64

```

from above we can see there is null values in genres and director columns

Next, we will deal with the features with missing values in genres and director I registered in tmdb official api read its docs, got an api key to use it in my requests

```
[ ]: # in this cell will deal with the missing genres values
```

```

# here is the api key that i got from tmdb api
api_key = "569f4c28a6676fc8b0c55407a23562a3"
# making a dataframe of the missing genres.

```

```

nan_genres = df.query('genres.isnull()')

# here we define a function to get the missing genre

def get_genre(movie_id):
    """
    this function takes the movie id as an argument,
    make the requests to tmdb api,
    then return the genres
    """
    genres = []
    resp = requests.get(
        "https://api.themoviedb.org/3/movie/{}?
    ↪api_key={}&language=en-US&append_to_response=credits".format(movie_id,
    ↪api_key))
    resp = json.loads(resp.content)
    try:
        # Get genres
        for genre in resp["genres"]:
            genres.append(str(genre["name"]).strip())
        genres = "|".join(genres)
    except:
        genres = 'other'
    return genres

# here we use the get_genre function to update the null values
for tmdb_id in nan_genres.id:
    df.loc[df.id == tmdb_id, "genres"] = get_genre(tmdb_id)

```

```

[ ]: # making a dataframe of the missing genres.
nan_director = df.query('director.isnull()')

# defining a function to grab the missing directors.

def get_director(movie_id):
    """
    this function takes the movie id as an argument,
    make the requests to tmdb api,
    then return the director
    """
    director = ""
    resp = requests.get(

```

```

        "https://api.themoviedb.org/3/movie/{}?
↳api_key={}&language=en-US&append_to_response=credits".format(movie_id,
↳api_key))
    resp = json.loads(resp.content)
    try:
        # Get director name
        for p in resp["credits"]["crew"]:
            if p["job"] == "Director":
                director = p["name"].strip()
    except:
        director = 'other'
    return director

# here we use the get_director function to update the null values
for tmdb_id in nan_director.id:
    df.loc[df.id == tmdb_id, "director"] = get_director(tmdb_id)

```

```

[ ]: # lets check again if there is still null values
df[['director', 'genres']].isnull().sum()

```

```

[ ]: director    0
     genres      0
     dtype: int64

```

next we will deal with zero values in the runtime, budget and revenue columns

```

[ ]: # making a dataframe of the rows of zero runtime values.
zero_runtime_df = df.query('runtime == 0')

def get_runtime(movie_id):
    """
    this function takes the movie id as an argument,
    make the requests to tmdb api,
    then return the runtime
    """
    resp = requests.get(
        "https://api.themoviedb.org/3/movie/{}?api_key={}&language=en-US".
↳format(movie_id, api_key))
    resp = json.loads(resp.content)

    try:
        runtime = resp['runtime']
    except:
        runtime = 0
    return runtime

```

```

# here we use the get_runtime function to update the zero values
runtime_count = 0 # this is counter of how many values has been updated.
for tmdb_id in zero_runtime_df.id:
    runtime = get_runtime(tmdb_id)
    if runtime != 0:
        df.loc[df.id == tmdb_id, 'runtime'] = runtime
        runtime_count += 1
print('There are {} runtime values has been updated'.format(runtime_count))

```

There are 28 runtime values has been updated

```

[ ]: # making a dataframe of the rows of zero budget values.
zero_budget_df = df.query('budget == 0')

def get_budget(movie_id):
    """
    this function takes the movie id as an argument,
    make the requests to tmdb api,
    then return the budget
    """
    resp = requests.get(
        "https://api.themoviedb.org/3/movie/{}?api_key={}&language=en-US".
        format(movie_id, api_key))
    resp = json.loads(resp.content)

    try:
        budget = resp['budget']
    except:
        budget = 0
    return budget

# here we use the get_budget function to update the zero values
budget_count = 0 # this is counter of how many values has been updated.
for tmdb_id in zero_budget_df.id:
    budget = get_budget(tmdb_id)
    if budget != 0:
        df.loc[df.id == tmdb_id, 'budget'] = budget
        budget_count += 1
print('There are {} budget values has been updated'.format(budget_count))

```

There are 1017 budget values has been updated

```

[ ]: # making a dataframe of the rows of zero revenue values.
zero_revenue_df = df.query('revenue == 0')

```



```

def get_revenue(movie_id):
    """
    this function takes the movie id as an argument,
    make the requests to tmdb api,
    then return the revenue
    """
    resp = requests.get(
        "https://api.themoviedb.org/3/movie/{}?api_key={}&language=en-US".
        format(movie_id, api_key))
    resp = json.loads(resp.content)

    try:
        revenue = resp['revenue']
    except:
        revenue = 0
    return revenue

# here we use the get_revenue function to update the zero values
revenue_count = 0 # this is counter of how many values has been updated.
for tmdb_id in zero_revenue_df.id:
    revenue = get_revenue(tmdb_id)
    if revenue != 0:
        df.loc[df.id == tmdb_id, 'revenue'] = revenue
        revenue_count += 1
print('There are {} revenue values has been updated'.format(revenue_count))

```

There are 1266 revenue values has been updated

now check how many zeros again

```

[ ]: print('zero values count in budget column was [5696] and now is: ',
        (df['budget'] == 0).sum())
print('zero values count in revenue column was [6016] and now is: ',
        (df['revenue'] == 0).sum())
print('zero values count in runtime column was [31] and now is: ',
        (df['runtime'] == 0).sum())

```

```

zero values count in budget column was [5696] and now is: 4679
zero values count in revenue column was [6016] and now is: 4750
zero values count in runtime column was [31] and now is: 3

```

> there is still featur with zero valus. > regarding runtime there is only still 3 movies with 0 runtime, we can get them manually using online search, i used (<https://en.wikipedia.org/>) > regarding budget and revenue, we will replace zeros with the median value

```

[ ]: df.query('runtime == 0')

```

```
[ ]:      id  popularity  budget  revenue  original_title  director \
616   325843    0.026797      0        0    the outfield michael goldfine
2370  127717    0.081892      0        0  freshman father michael scott
5992  173847    0.096455      0        0  amiche da morire giorgia farina

      runtime      genres  vote_count  vote_average  release_year
616         0  drama|comedy         20          6.6         2015
2370         0  Drama|TV Movie         12          5.8         2010
5992         0  romance|crime|comedy        30          5.5         2013
```

```
[ ]: # after online search, we will set the runtime values we got.
#
df.loc[df.id == 325843, 'runtime'] = 90
df.loc[df.id == 127717, 'runtime'] = 87
df.loc[df.id == 173847, 'runtime'] = 103

# check for 0 runtime again
df[df['runtime']==0].runtime.sum()
```

```
[ ]: 0
```

```
[ ]: # replacing the 0 budget and 0 revenue with the median value
df.budget.replace(to_replace=0, value=df.budget.median(), inplace=True)
df.revenue.replace(to_replace=0, value=df.revenue.median(), inplace=True)
```

```
[ ]: print('zero values count in budget column was [5696] and now is: ',
      ↪(df['budget'] == 0).sum())
print('zero values count in revenue column was [6016] and now is: ',
      ↪(df['revenue'] == 0).sum())
print('zero values count in runtime column was [31] and now is: ',
      ↪(df['runtime'] == 0).sum())
```

```
zero values count in budget column was [5696] and now is:  0
zero values count in revenue column was [6016] and now is:  0
zero values count in runtime column was [31] and now is:  0
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 10865 entries, 0 to 10865
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           10865 non-null  int64
1   popularity   10865 non-null  float64
2   budget       10865 non-null  int64
3   revenue      10865 non-null  int64
```

```

4  original_title  10865 non-null  object
5  director       10865 non-null  object
6  runtime        10865 non-null  int64
7  genres         10865 non-null  object
8  vote_count     10865 non-null  int64
9  vote_average   10865 non-null  float64
10 release_year   10865 non-null  int64
dtypes: float64(2), int64(6), object(3)
memory usage: 1018.6+ KB

```

```
[ ]: # saving the cleaned data into new csv file
df.to_csv('tmdb-cleaned.csv',index=False)
```

```
[ ]: cleaned_df = pd.read_csv('tmdb-cleaned.csv')
cleaned_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10865 entries, 0 to 10864
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10865 non-null  int64
1   popularity            10865 non-null  float64
2   budget               10865 non-null  int64
3   revenue              10865 non-null  int64
4   original_title        10865 non-null  object
5   director              10865 non-null  object
6   runtime              10865 non-null  int64
7   genres               10865 non-null  object
8   vote_count           10865 non-null  int64
9   vote_average          10865 non-null  float64
10  release_year          10865 non-null  int64
dtypes: float64(2), int64(6), object(3)
memory usage: 933.8+ KB

```

after reading the cleaned df i found there is 5 missing directors, i got them using online search and will update them manually

```
[ ]: cleaned_df[cleaned_df.director.isnull()]
```

```

[ ]: Empty DataFrame
Columns: [id, popularity, budget, revenue, original_title, director, runtime,
genres, vote_count, vote_average, release_year]
Index: []

```

i used (<https://en.wikipedia.org/>) to get the missing directors and assigned them manually as follows

```
[ ]: cleaned_df.loc[cleaned_df.id == 259910,'director'] = 'brad baruh'
cleaned_df.loc[cleaned_df.id == 253675,'director'] = 'susannah ward'
cleaned_df.loc[cleaned_df.id == 301235,'director'] = 'phil churchward'
cleaned_df.loc[cleaned_df.id == 17171,'director'] = "andrew dominik"
cleaned_df.loc[cleaned_df.id == 46188,'director'] = 'david cherkasskiy'
```


```
[ ]: # saving to csv file again
df.to_csv('tmdb-cleaned.csv',index=False)
```

## Exploratory Data Analysis

**1.0.3 Q1:** Lets see how the number of produced movies changes over the years?

```
[ ]: cleaned_df = pd.read_csv('tmdb-cleaned.csv')
```

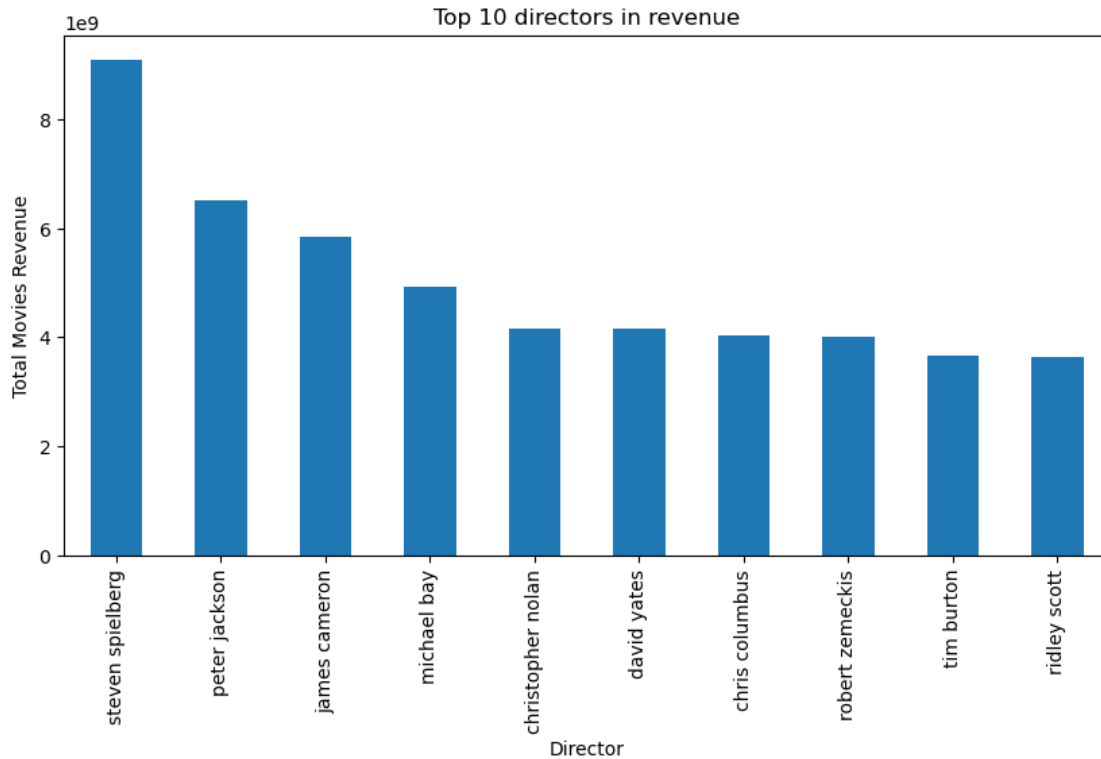
```
[ ]: cleaned_df.groupby('release_year')['original_title'].count().plot(kind='bar',
↪figsize=(12,5));
```



As we see from the above chart, No. of produced movies starts increasing dramatically in 1993, and keep increasing over the years to reach the top in 2014

**1.0.4 Q2:** Lets see Top 10 directors have the most revenue for thier films?


```
[ ]: cleaned_df.groupby('director')['revenue'].sum().nlargest(10).plot(kind='bar',
↪x='director', y='revenue',xlabel='Director', ylabel='Total Movies Revenue',
↪figsize=(10,5), title='Top 10 directors in revenue');
```



from the above chart, we see that the top 10 directors whose movies have the highest revenue:  
 1- steven spielberg 2-peter jackson 3-james cameron 4-michael bay 5-christopher nolan 6-david yates  
 7-chris columbus 8-robert zemeckis 9-tim burton 10-ridley scott

### 1.0.5 Q3: what is the top 10 movies in revenue?

```
[ ]: cleaned_df.nlargest(10, 'revenue')[['original_title', 'revenue']].
      plot(kind='bar', x='original_title', y='revenue', xlabel='Movie',
      ylabel='Revenue', title='Top 10 Movies in Revenue', figsize=(10,5));
```




the chart shows that the highest movie in revenue is: 1-avatar 2-star wars: the force awakens 3-titanic 4-the avengers 5-jurassic world 6-furious 7 7-avengers: age of ultron 8-harry potter and the deathly hallows: part2 9-frozen 10-iron man 3

#### 1.0.6 Q4: what is the most produced genres ?

```
[ ]: genres_df= cleaned_df.genres.str.split('|', expand=True)
genres_df = pd.DataFrame(genres_df.stack())
genres_df.rename(columns={0: 'genres_adj'}, inplace=True)
genres_df.genres_adj =genres_df.genres_adj.str.lower()
```


```
[ ]: genres_df.genres_adj.value_counts().plot(kind='bar', figsize=(10,5));
```



as we see the most produced genre is drama, then comedy, then thriller and action

### 1.0.7 Q5: does popularity affect the revenue?

```
[ ]: plt.scatter(x=cleaned_df.popularity, y=cleaned_df.revenue, marker='.', linewidths=1, );
plt.xticks(range(0,1+int(cleaned_df.popularity.max()),2 ));
plt.xlabel('Popularity');
plt.ylabel('Revenue');
plt.title('Populatiy Vs. Revenue');
```




there is a positive correlation between popularity and revenue

#### 1.0.8 Q6: does runtime affect the revenue?

```
[ ]: plt.scatter(x=cleaned_df.runtime, y=cleaned_df.revenue, marker='.',  
               ↪linewidths=1, );  
# plt.xticks(range(0,1+int(cleaned_df.popularity.max()),2 ));  
plt.xlabel('Runtime');  
plt.ylabel('Revenue');  
plt.title('Runtime Vs. Revenue');
```






there is a weak positive correlation between runtime and revenue

#### 1.0.9 Q7: does vote count affect the revenue?


```
[ ]: plt.scatter(x=cleaned_df.vote_count, y=cleaned_df.revenue, marker='.', linewidths=1);
plt.xlabel('Vote Count');
plt.ylabel('Revenue');
plt.title('Vote Count Vs. Revenue');
```



there is a positive correlation between Vote Count and Revenue

#### 1.0.10 Q8: does vote average affect the revenue?

```
[ ]: plt.scatter(x=cleaned_df.vote_average, y=cleaned_df.revenue, marker='.', linewidths=1);
plt.xlabel('Vote Average');
plt.ylabel('Revenue');
plt.title('Vote Average Vs. Revenue');
```



there is a positive correlation between Vote Average and Revenue

## 2 conclusion

> steven spielberg is the director with highest revenue for his total moveies till 2015 > the most produced movies per year was 700 in 2014, and the minimum is 31 in 1961

> till 2015 avatar had the highest revenue at all with 2781505847 > most produced genres are drama, then comedy, then thriller and action recpevly > there is a weake positive correlation between revenue and (runtime, popularity, vote count and vote average)

### 2.0.1 limitations:

> there is many zero values in budget and revenue features we replaced it with the median, wich could lead to inaccuracy > There is no currency unit for budget and revenue features so they may be in different currencies for different countries.